

Code Transcript

Q1.

```
import pandas as pd
df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')

df.sort_values(by='subscribers', ascending=False)
print(df.head(10))
```

Output:

rank	Youtuber	subscribers	video views	category	...	Population	Unemployment rate	Urban_population	Latitude	Longitude
0	1 T-Series	245000000.0	2.280000e+11	Music	...	1.366418e+09	5.36	471031528.0	20.593684	78.962880
1	2 YouTube Movies	170000000.0	0.000000e+00	Film & Animation	...	3.282395e+08	14.70	270663028.0	37.090240	-95.712891
2	3 MrBeast	166000000.0	2.836884e+10	Entertainment	...	3.282395e+08	14.70	270663028.0	37.090240	-95.712891
3	4 Cocomelon - Nursery Rhymes	162000000.0	1.640000e+11	Education	...	3.282395e+08	14.70	270663028.0	37.090240	-95.712891
4	5 SET India	159000000.0	1.480000e+11	Shows	...	1.366418e+09	5.36	471031528.0	20.593684	78.962880
5	6 Music	119000000.0	0.000000e+00	NaN	...	NaN	NaN	NaN	NaN	NaN
6	7 yyy Kids Diana Show	112000000.0	9.324704e+10	People & Blogs	...	3.282395e+08	14.70	270663028.0	37.090240	-95.712891
7	8 PewDiePie	111000000.0	2.905804e+10	Gaming	...	1.262266e+08	2.29	115782416.0	36.204824	138.252924
8	9 Like Nastya	106000000.0	9.047906e+10	People & Blogs	...	1.443735e+08	4.59	107683889.0	61.524010	105.318756
9	10 vLad and Niki	98900000.0	7.718017e+10	Entertainment	...	3.282395e+08	14.70	270663028.0	37.090240	-95.712891

Q2.

```
import pandas as pd
df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
category_avg = df.groupby('category')['subscribers'].mean()
category_avg_sorted = category_avg.sort_values(ascending=False)

print("Category with the highest average subscribers:")
print(category_avg_sorted.head(1))
```

Output:

```
Category with the highest average subscribers:
category
Shows      4.161538e+07
```

Q3.

```
import pandas as pd

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
avg_uploads = df.groupby('category')['uploads'].mean().round(0).astype(int)
avg_uploads_sorted = avg_uploads.sort_values(ascending=False)

print("Average number of uploads per category (rounded):")
print(avg_uploads_sorted)
```

Output:

```
Average number of uploads per category
category
News & Politics      112484
Nonprofits & Activism 102912
Shows                 27444
Sports                19130
Entertainment         12052
People & Blogs        9257
Trailers               6839
Gaming                4285
Pets & Animals         3563
Movies                 3553
Education              3087
Film & Animation       2862
Music                  2326
Science & Technology   2114
Howto & Style           1696
Autos & Vehicles        1551
Comedy                 1203
Travel & Events          766
```

Q4.

```
import pandas as pd

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
top_countries = df['Country'].value_counts().head()

print("Top 5 countries with the highest number of YouTube channels:")
print(top_countries)
```

Output:

Top 5 countries with the highest number of YouTube channels:

Country	
United States	315
India	169
Brazil	62
United Kingdom	44
Mexico	33

Q5.

```
import pandas as pd
import matplotlib.pyplot as plt


df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
distribution = pd.crosstab(df['category'], df['channel_type'])

plt.figure(figsize=(14, 7))
distribution.plot(kind='bar', stacked=True, colormap='Set3')

plt.title('Distribution of Channel Types Across Categories')
plt.xlabel('Category')
plt.ylabel('Number of Channels')
plt.xticks(rotation=45, ha='right')
plt.legend(title='Channel Type', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()

plt.show()
```

Output:



Q6.

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')

df_clean = df.dropna(subset=['subscribers', 'video views'])
df_clean = df_clean[(df_clean['subscribers'] > 0) & (df_clean['video views'] > 0)]

correlation = df_clean['subscribers'].corr(df_clean['video views'])
print(f"Correlation between subscribers and video views: {correlation:.2f}")

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_clean, x='subscribers', y='video views', alpha=0.6, color='purple', edgecolor=None)

plt.xscale('log')
plt.yscale('log')


plt.title('Correlation Between Subscribers and Video Views')
plt.xlabel('Subscribers (log scale)')
plt.ylabel('Video Views (log scale)')
plt.grid(True)
plt.tight_layout()

plt.show()
```

Output:

Strong Positive Correlation observed

Correlation between subscribers and video views: 0.83



Q7.


```
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df['lowest_monthly_earnings'] = pd.to_numeric(df['lowest_monthly_earnings'], errors='coerce')
df['highest_monthly_earnings'] = pd.to_numeric(df['highest_monthly_earnings'], errors='coerce')
df_clean = df.dropna(subset=['category', 'lowest_monthly_earnings', 'highest_monthly_earnings'])

monthly_earnings = df_clean.groupby('category')[['lowest_monthly_earnings', 'highest_monthly_earnings']].mean()
monthly_earnings = (monthly_earnings / 1000).round(1) # Divide by 1000 and round to 1 decimal place
monthly_earnings_sorted = monthly_earnings.sort_values(by='highest_monthly_earnings', ascending=False)
monthly_earnings_sorted.plot(kind='bar', figsize=(12, 6), color=['skyblue', 'orange'])

plt.title('Average Monthly Earnings by Category (in Thousands)')
plt.xlabel('Category')
plt.ylabel('Earnings (in $1,000s)')
plt.xticks(rotation=45, ha='right')
plt.legend(['Lowest Monthly Earnings', 'Highest Monthly Earnings'])
plt.grid(True, axis='y')
plt.tight_layout()
plt.show()
```

Output:



Q8.

```
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df['subscribers_for_last_30_days'] = pd.to_numeric(df['subscribers_for_last_30_days'], errors='coerce')
df_clean = df.dropna(subset=['category', 'subscribers_for_last_30_days'])


sub_trend = df_clean.groupby('category')['subscribers_for_last_30_days'].mean()
sub_trend = (sub_trend / 1000).round(1)

sub_trend_sorted = sub_trend.sort_values(ascending=False)

plt.figure(figsize=(12, 6))
sub_trend_sorted.plot(kind='bar', color='mediumseagreen')

plt.title('Average Subscribers Gained in Last 30 Days by Category (in Thousands)')
plt.xlabel('Category')
plt.ylabel('Subscribers Gained (in 1,000s)')
plt.xticks(rotation=45, ha='right')
plt.grid(True, axis='y')
plt.tight_layout()
plt.show()
```

Output:



Q9.

```
import pandas as pd

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df['highest_yearly_earnings'] = pd.to_numeric(df['highest_yearly_earnings'], errors='coerce')
df_clean = df.dropna(subset=['highest_yearly_earnings'])

Q1 = df_clean['highest_yearly_earnings'].quantile(0.25)
Q3 = df_clean['highest_yearly_earnings'].quantile(0.75)
IQR = Q3 - Q1
upper_bound = Q3 + 1.5 * IQR
outliers = df_clean[df_clean['highest_yearly_earnings'] > upper_bound]

print(f"Upper bound for outliers: {round(upper_bound):,} USD\n")
print("YouTube Channels with Outlier Yearly Earnings:\n")
print(outliers[['Youtuber', 'category', 'highest_yearly_earnings']].sort_values(by='highest_yearly_earnings', ascending=False))
```

Output:

Upper bound for outliers: 17,468,875 USD
 YouTube Channels with Outlier Yearly Earnings:

	Youtuber	category	highest_yearly_earnings
495	ýýýýýýýýýý KIMPRO	NaN	163400000.0
417	DaFuq! ?Boom!	Film & Animation	110600000.0
0	T-Series	Music	108400000.0
0	T-Series	Music	108400000.0
0	T-Series	Music	108400000.0
302	KL BRO Biju Rithvik	Entertainment	97600000.0
0	T-Series	Music	108400000.0
0	T-Series	Music	108400000.0
302	KL BRO Biju Rithvik	Entertainment	97600000.0
3	Cocomelon - Nursery Rhymes	Education	94800000.0
..
134	The Weeknd	Music	17900000.0
502	News 24	News & Politics	17800000.0
779	SEVENGERS	Comedy	17700000.0
908	Susy Mouriz	Entertainment	17700000.0
283	Tilak	Film & Animation	17600000.0

Q10.


```
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df['created_year'] = pd.to_numeric(df['created_year'], errors='coerce')
df = df.dropna(subset=['created_year'])
channel_trend = df['created_year'].value_counts().sort_index()

plt.figure(figsize=(12, 6))
plt.plot(channel_trend.index, channel_trend.values, marker='o', linestyle='-', color='teal')

plt.title('Trend of YouTube Channel Creations Over the Years')
plt.xlabel('Year')
plt.ylabel('Number of Channels Created')
plt.grid(True)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

Output:



Q11.

```

import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df_clean = df.dropna(subset=['Country', 'Gross tertiary education enrollment (%)'])
channel_counts = df_clean['Country'].value_counts().rename_axis('Country').reset_index(name='Channel_Count')
edu_avg = df_clean.groupby('Country')['Gross tertiary education enrollment (%)'].mean().reset_index()
merged = pd.merge(channel_counts, edu_avg, on='Country')

plt.figure(figsize=(10, 6))
plt.scatter(merged['Gross tertiary education enrollment (%)'], merged['Channel_Count'], color='purple')
plt.title('Tertiary Education Enrollment vs YouTube Channel Count per Country')
plt.xlabel('Gross Tertiary Education Enrollment (%)')
plt.ylabel('Number of YouTube Channels')
plt.grid(True)
plt.tight_layout()
plt.show()


correlation = merged['Gross tertiary education enrollment (%)'].corr(merged['Channel_Count'])
print(f"Correlation coefficient: {correlation:.2f}")

```

Output:

No meaningful relationship observed, only a few anomalous plots seen on the scatter plot.

Correlation coefficient: 0.11



Q12.

```

import pandas as pd
import matplotlib.pyplot as plt


df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df_clean = df.dropna(subset=['Country', 'Unemployment rate'])
channel_counts = df_clean['Country'].value_counts().head(10)
top_countries = channel_counts.index.tolist()
top_df = df_clean[df_clean['Country'].isin(top_countries)]
avg_unemployment = top_df.groupby('Country')['Unemployment rate'].mean().sort_values(ascending=False)

plt.figure(figsize=(12, 6))
avg_unemployment.plot(kind='bar', color='coral')

plt.title('Average Unemployment Rate in Top 10 Countries by YouTube Channel Count')
plt.ylabel('Unemployment Rate (%)')
plt.xlabel('Country')
plt.grid(axis='y', linestyle='--', linewidth=0.5)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

Output:




Q13.

```
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df_clean = df.dropna(subset=['Country', 'Urban_population'])
df_clean['Country'] = df_clean['Country'].str.strip().str.title()
urban_pop_by_country = (
    df_clean.groupby('Country')['Urban_population']
    .mean()
    .round(2)
    .sort_values(ascending=True)
)
plt.figure(figsize=(12, 20))
urban_pop_by_country.plot(kind='barh', color='skyblue')

plt.title('Average Urban Population % by Country with YouTube Channels')
plt.xlabel('Urban Population (%)')
plt.ylabel('Country')
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

Output:



Q14.


```
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df_clean = df.dropna(subset=['Latitude', 'Longitude'])

plt.figure(figsize=(12, 6))
plt.scatter(df_clean['Longitude'], df_clean['Latitude'], alpha=0.5, s=30, color='purple')

plt.title('Geographic Distribution of YouTube Channels')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

Output:



Q15.

```

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import linregress

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df_clean = df.dropna(subset=['subscribers', 'Population', 'Country'])
country_group = df_clean.groupby('Country').agg({
    'subscribers': 'sum',
    'Population': 'mean'
}).reset_index()
country_group = country_group[(country_group['subscribers'] > 0) & (country_group['Population'] > 0)]
country_group['log_subs'] = np.log10(country_group['subscribers'])
country_group['log_pop'] = np.log10(country_group['Population'])

slope, intercept, r_value, _, _ = linregress(country_group['log_pop'], country_group['log_subs'])


plt.figure(figsize=(10, 6))
plt.scatter(country_group['log_pop'], country_group['log_subs'], color='steelblue', alpha=0.7, label='Country Data')

x_vals = np.array(plt.gca().get_xlim())
y_vals = intercept + slope * x_vals
plt.plot(x_vals, y_vals, '--', color='darkred', label=f'Trendline (r = {r_value:.2f})')

plt.xlabel('Log10 of Country Population')
plt.ylabel('Log10 of Total Subscribers')
plt.title('Correlation Between Country Population and YouTube Subscribers\n(Log Scaled for Better Visualization)')
plt.grid(True, linestyle='--', alpha=0.6)
plt.legend()
plt.tight_layout()
plt.show()

```

Output: There is a moderate positive correlation with a value of $r=0.51$ as shown in the graph.



Q16.

```

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df_clean = df.dropna(subset=['Country', 'Population'])
channel_counts = df_clean['Country'].value_counts().head(10)
population_data = df_clean[df_clean['Country'].isin(channel_counts.index)][['Country', 'Population']]
population_data = population_data.drop_duplicates(subset='Country')

merged = pd.DataFrame({
    'Country': channel_counts.index,
    'YouTube Channels': channel_counts.values
})
merged = pd.merge(merged, population_data, on='Country')
merged['Log Population'] = np.log10(merged['Population'])

fig, ax1 = plt.subplots(figsize=(12, 6))


ax1.bar(merged['Country'], merged['YouTube Channels'], color='steelblue', label='YouTube Channels')
ax1.set_xlabel('Country')
ax1.set_ylabel('YouTube Channels', color='steelblue')
ax1.tick_params(axis='y', labelcolor='steelblue')
ax1.set_xticklabels(merged['Country'], rotation=45)

ax2 = ax1.twinx()
ax2.plot(merged['Country'], merged['Log Population'], color='darkorange', marker='o', linewidth=2, label='Log10(Population)')
ax2.set_xlabel('Log10 of Population', color='darkorange')
ax2.tick_params(axis='y', labelcolor='darkorange')

plt.title('Top 10 Countries by YouTube Channels vs Log10(Population)')
plt.tight_layout()
plt.grid(axis='y', linestyle='--', alpha=0.5)
plt.show()

```

Output:



Q17.


```
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import pearsonr

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df_clean = df[['Country', 'subscribers_for_last_30_days', 'Unemployment rate']].dropna()
df_clean['subscribers_for_last_30_days'] = pd.to_numeric(df_clean['subscribers_for_last_30_days'], errors='coerce')
df_clean['Unemployment rate'] = pd.to_numeric(df_clean['Unemployment rate'], errors='coerce')
df_clean = df_clean.dropna()
grouped = df_clean.groupby('Country').agg({
    'subscribers_for_last_30_days': 'mean',
    'Unemployment rate': 'mean'
}).reset_index()

corr, _ = pearsonr(grouped['subscribers_for_last_30_days'], grouped['Unemployment rate'])
print(f"Pearson correlation: {corr:.2f}")

plt.figure(figsize=(10,6))
plt.scatter(grouped['Unemployment rate'], grouped['subscribers_for_last_30_days'] / 100000, color='mediumseagreen')
plt.title('Correlation Between Unemployment Rate and Subscribers Gained (Last 30 Days)')
plt.xlabel('Unemployment Rate (%)')
plt.ylabel('Avg Subscribers Gained (in 100k)')
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

Output: Pearson correlation value is -0.06 which indicates no correlation.



Q18.

```

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df_clean = df[['channel_type', 'video_views_for_the_last_30_days']].dropna()
df_clean['video_views_for_the_last_30_days'] = pd.to_numeric(
    df_clean['video_views_for_the_last_30_days'], errors='coerce'
)
df_clean = df_clean.dropna()
df_clean = df_clean[df_clean['video_views_for_the_last_30_days'] > 0]

grouped = df_clean.groupby('channel_type').agg({
    'video_views_for_the_last_30_days': 'mean',
    'channel_type': 'count'
}).rename(columns={
    'video_views_for_the_last_30_days': 'Average Views (Last 30 Days)',
    'channel_type': 'Channel Count'
}).reset_index()
grouped = grouped.sort_values('Average Views (Last 30 Days)', ascending=False)
fig, ax1 = plt.subplots(figsize=(14, 6))


ax1.bar(grouped['channel_type'], grouped['Channel Count'], color='skyblue', label='Channel Count')
ax1.set_xlabel('Channel Type')
ax1.set_ylabel('Channel Count', color='skyblue')
ax1.tick_params(axis='y', labelcolor='skyblue')
ax1.set_xticklabels(grouped['channel_type'], rotation=45, ha='right')

ax2 = ax1.twinx()
ax2.plot(grouped['channel_type'], grouped['Average Views (Last 30 Days)'] / 1e6, color='darkorange', marker='o', label='Avg Views (in Millions)')
ax2.set_ylabel('Average Views (Last 30 Days) [in Millions]', color='darkorange')
ax2.tick_params(axis='y', labelcolor='darkorange')

plt.title('Channel Type: Count vs Avg Views (Last 30 Days)')
plt.grid(axis='y', linestyle='--', alpha=0.4)
fig.tight_layout()
plt.show()

```

Output:



Q19.

Output:

Based on the available dataset, there is no timestamped information about individual video uploads because only the total number of uploads per channel and their creation month are provided.

Therefore, we do not have sufficient data to accurately determine seasonal trends in upload behavior over time.

Q20.

```
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib.ticker import FuncFormatter
from datetime import datetime

df = pd.read_csv("Global YouTube Statistics.csv", encoding='latin1')
df.columns = df.columns.str.strip()
df['subscribers'] = pd.to_numeric(df['subscribers'], errors='coerce')
df['created_year'] = pd.to_numeric(df['created_year'], errors='coerce')
df = df.dropna(subset=['subscribers', 'created_year'])
df = df[(df['created_year'] >= 2006) & (df['subscribers'] > 0)]

current_year = datetime.now().year
df['years_since_creation'] = current_year - df['created_year']
df = df[df['years_since_creation'] > 0]
df['months_since_creation'] = df['years_since_creation'] * 12
df['subs_per_month'] = df['subscribers'] / df['months_since_creation']
overall_avg = df['subs_per_month'].mean()

plt.figure(figsize=(12, 6))
plt.hist(df['subs_per_month'], bins=100, color="#4C72B0", edgecolor='black', log=True)
plt.gca().xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{int(x/1000)}K'))
plt.axvline(overall_avg, color='crimson', linestyle='--', linewidth=2, label=f'Avg ≈ {overall_avg:.0f}')

plt.title('Subscribers Gained Per Month(approx value only)', fontsize=14, fontweight='bold')
plt.xlabel('Subscribers per Month', fontsize=12)
plt.ylabel('Number of Channels (log scale)', fontsize=12)
plt.legend()
plt.grid(True, linestyle='--', alpha=0.4)
plt.tight_layout()
plt.show()
```

Output:

